

# Software Quality for the Semantic Web


Eduard Kamburjan

*Based on work with many collaborators: Tobias John, Einar Broch Johnsen, Dominic Steinhöfel, David Chaves Fraga, Romana Pernisch, Oscar Corcho, ...*


onto:NEXUS Workshop 05.10.2025

IT UNIVERSITY OF COPENHAGEN

# What is a Knowledge Graph, and why do we care?



# What is a Knowledge Graph, and why do we care?




## The Promise of Knowledge Graphs

A Knowledge Graph is a graph that provides high-quality semantic data to users.

- Neuro-symbolic AI: context and taming hallucinations in GenAI, ...
- Data integration: Data of high quality with agreed upon semantics, ...
- Engineering: Connecting system models with data, ...

# What is a Knowledge Graph, and why do we care?



## The Promise of Knowledge Graphs

A Knowledge Graph is a graph that provides high-quality semantic data to users.

- Neuro-symbolic AI: context and taming hallucinations in GenAI, ...
- Data integration: Data of high quality with agreed upon semantics, ...
- Engineering: Connecting system models with data, ...

If the central promise is quality, how do we ensure it?

# What is a Knowledge Graph really?

Data Quality

Software Quality

# What is a Knowledge Graph really?

## Data Quality

## Software Quality

- Data engineering pipelines with numerous tools on general data quality
- Graph specific technologies: SHACL shapes, SPARQL queries as constraints
- Formal semantics and reasoners
- Lot of different methodologies

# What is a Knowledge Graph really?

## Data Quality

- Data engineering pipelines with numerous tools on general data quality
- Graph specific technologies: SHACL shapes, SPARQL queries as constraints
- Formal semantics and reasoners
- Lot of different methodologies

## Software Quality

- What about all these tools?
- What about all these pipelines?



## Software Turtles all the way down

### Knowledge Graphs

- A KG is a data set, generated by a set of interacting software components.
- The quality of the KG is determined also by their software quality.

# Software Turtles all the way down

## Knowledge Graphs

- A KG is a data set, generated by a set of interacting software components.
- The quality of the KG is determined also by their software quality.

## Software Quality is Important

- Five papers in high-impact venues (including nature) retracted after a bug in python implementation of analysis algorithm  
*[Miller, Software problem leads to five retractions., 2007]*
- Faulty analysis leads to wrong data basis for decision about austerity in Europe  
*[Herndon et al., Does High Public Debt Consistently Stifle Economic Growth? A Critique of Reinhart and Rogoff, 2013]*
- “Replication crisis” w.r.t. Jupyter notebooks: less than 25% are runnable  
*[Pimentel et al., Understanding and improving the quality and reproducibility of Jupyter notebooks, 2021]*

# Agenda

- Testing software for knowledge graphs (x2)
- Dependency analysis for knowledge graph construction

# A very short primer on knowledge graphs

## Triple-Based Knowledge Representation

*Knowledge Graphs* are a framework to (a) represent, (b) reason over, and (c) query domain knowledge and data.

# A very short primer on knowledge graphs

## Triple-Based Knowledge Representation

*Knowledge Graphs* are a framework to (a) represent, (b) reason over, and (c) query domain knowledge and data.

## W3C Standards

RDF for data, OWL for knowledge, SPARQL for queries.

# A very short primer on knowledge graphs

## Triple-Based Knowledge Representation

*Knowledge Graphs* are a framework to (a) represent, (b) reason over, and (c) query domain knowledge and data.

## W3C Standards

RDF for data, OWL for knowledge, SPARQL for queries.

RDF:  
Paul a Person. Peter a Person. Maria a Person.  
Paul hasChild Peter. Peter hasChild Maria.

# A very short primer on knowledge graphs

## Triple-Based Knowledge Representation

*Knowledge Graphs* are a framework to (a) represent, (b) reason over, and (c) query domain knowledge and data.

## W3C Standards

RDF for data, OWL for knowledge, SPARQL for queries.

RDF:  
Paul a Person. Peter a Person. Maria a Person.  
Paul hasChild Peter. Peter hasChild Maria.

OWL:  
hasChild some (hasChild some Person) subClassOf GrandParent  
 $\exists \text{hasChild. } \exists \text{hasChild. Person} \sqsubseteq \text{GrandParent}$

# A very short primer on knowledge graphs

## Triple-Based Knowledge Representation

*Knowledge Graphs* are a framework to (a) represent, (b) reason over, and (c) query domain knowledge and data.

## W3C Standards

RDF for data, OWL for knowledge, SPARQL for queries.

RDF:  
Paul a Person. Peter a Person. Maria a Person.  
Paul hasChild Peter. Peter hasChild Maria.

OWL:  
hasChild some (hasChild some Person) subClassOf GrandParent  
 $\exists \text{hasChild. } \exists \text{hasChild. Person} \sqsubseteq \text{GrandParent}$

SPARQL: SELECT ?x WHERE { ?x a GrandParent }

# Testing (I): Language-based Fuzzing

---



# Automated Testing

## Problems

- How can we automatically test the general purpose tools in OE?
- How can we test the integration of ontologies with other software?
- How can we specify the integration of ontologies with other software?

- Solver and database engines are hard to test in general
- RDF has per se very little structure to constraint input generation

# Automated Testing

- Language-based approach to generate random graphs and ontologies with ISLa
- Two grammars: RDF/TTL and OWL functional syntax

```
1 <ontology>      ::= "Ontology (" <declarations> " " <axioms> ")"
2 <axioms>         ::= <axiom> | <axiom> "\n" <axioms>
3 <axiom>          ::= <classAxiom> | <assertion> | <dataTypeDefinition> | [...]
4 [...]             ...
5 <literal>        ::= <typedLiteral> | <stringNoLang> | <stringWithLang>
6 <stringWithLang> ::= <QuotedString> <LanguageTag>
```

## Targets

- RDF/TTL parser and frontend utilities of Apache Jena and OWL-API
- Three OWL-EL reasoners via differential testing

## Automated Testing: Frontend Bugs

- First bug found in RDF 1.2 TTL standard with second generated file

```
<P:A> <B> <C>.  
@prefix P: <http://test.no#>
```
- Both parser have bugs in corner cases, despite a formal grammar in the standard!

```
<A> <B> -.7 . // fails to parse literal  
<A> <B> ; ; . // fails to parse double empty list
```
- OWL-API profile checker rejects all OWL-EL ontologies that use language tags

```
<A> <B> "test"^^xsd:String@dk_DK
```

# Automated Testing: OWL-EL Reasoners

## Test Targets

Three reasoners included by default with Protege


- HermiT (v.1.4.5.519)
- Pellet/Openllet (v.2.6.5)
- ELK (v.0.6.0)

## Test Procedure

- Generate new ontology, and ask all three reasoners if it is consistent and to derive all possible axioms
- If results are different (or exception is thrown), investigate
- Extra tool to reduce ontology by axiom pinpointing

# Automated Testing: OWL-EL Reasoners

- Found and reported 15 bugs, 13 from failed logical inference, 2 from exceptions
- Language tags and corner cases in the hierarchy



# Automated Testing: OWL-EL Reasoners

```
1 //ELK classifies as inconsistent
2 Prefix(:=<http://www.example.org/reasonerTester/>)
3 Ontology (
4     Declaration(Class(:B)) Declaration(Class(:A))
5     Declaration(DataProperty(:dr)) Declaration(NamedIndividual(:a))
6     EquivalentClasses( DataHasValue(:dr "s1"@fr) :A :B )
7     DisjointClasses( DataHasValue(:dr "s1"@en) :A )
8     ClassAssertion(:B :a))
```

```
1 //HermiT fails to derive DataPropertyAssertion(:dp :a "data")
2 Prefix(:=<http://www.example.org/reasonerTester/>)
3 Ontology (
4     Declaration(DataProperty(:dp)) Declaration(NamedIndividual(:a))
5     EquivalentClasses( ObjectOneOf(:a) DataHasValue(:dp "data") ))
```

## Automated Testing: Applicability

- Found numerous bugs in all tested tools, only with black box testing and limited tasks/oracles
- Proves that automated testing of general purpose tools for KGs is possible and feasible
- General purpose grammars a bit unhandy and need to be constraint by hand for more specific applications
- ISLa not optimal for high-volume generation


## **Testing (II): Mutation-based Integration Testing**

---




# Integration Testing

- Given a program, we often have an example KG it interacts with
- What exactly do we need to specify the program-KG interface?
- Mutation of KG to generate new inputs to program
- Challenges: Mutating KG depends on domain, program has implicit assumptions




# Integration Testing: Mutation Operators

- Prior work mutate single triples or axioms
- Too fine-grained for programs – removing one entity may change a whole sub graph



(c) Scenario with corresponding KG representation before mutating.



(d) Scenario with corresponding KG representation after mutating.

# Integration Testing: Robustness Mask

## Robustness Mask

- Not every consistent ontology is valid input
- Program has implicit assumptions about ontology
- Top-level ontology should probably not be mutated
- Additional SHACL shapes to constrain mutations

```
p := query(":isAt(:auv, ?p)")           AuvAtPipeline
inspect(p)                                a sh:NodeShape ;
S := query(":nextTo(p, ?s)")               sh:targetNode :auv ;
while S ≠ ∅ do                            sh:property [
    p := S.pop()                          sh:path :isAt ;
    if ¬inspected(p) then                sh:minCount 1 ;
        moveTo(p)                      sh:maxCount 1 ;
        inspect(p)                    sh:class :Pipe ;
        S := query(":nextTo(p, ?s)")   ] .
    end if
end while
```


## Domain-Specific Operators

- Defined per ontology or test suite
- Either directly implemented on KG (imported via Kotlin)
- Or by using SWRL syntax for rewriting


```
rdfmutate:newNode(?p) ∧ :Topping(?t) → :Pizza(?p) ∧ :hasTopping(?p, ?t)
```

- 59 relatively generic operators predefined
- Prototypical implementation based on rule-mining can automate initial domain-specific operators

# Integration Testing: Input Coverage



(a) Domain-specific mutation operators



(b) Learned operators

- Input feature coverage: How many features are used?
- Measured via OWL vocabulary
- Domain-specific operators can be used to force feature interactions

# Integration Testing: Results

## Targets

- SUAVE: Simulator for self-adaptive AUV based on ROS
- GeoSimulator: Simulator for geological process based on geological ontologies
- OWL-EL reasoners: Same setup

## Seed Ontologies

- Suave and GeoSimulator: Only one ontology as default example
- OWL-EL reasoners: 307 Ontologies from latest OWL reasoning competition

## Results

- SUAVE: Mistakes in OWL modeling
- GeoSimulator: No bugs
- OWL-EL reasoners: 6 additional bugs related to reasoning over class hierarchies

## Integration Testing: Conclusion

- Robustness mask useful for interface specification
- Even with automation, domain-specific operations require some work
- But easier to control and estimate compared to grammar-based fuzzing.
- Again, found bugs in non-trivial systems

# Dependency Analysis

---


ITU

# Dependencies for KGC


## Problem

Given an KGC pipeline, can we assess the impact of a change in a component?

- Impact analysis based on a dependency analysis
- Challenge: Some used language have no formal semantics
- Challenge: Notion of dependencies not used in KGC
- First study on dependencies for impact analysis and bug detection



# Dependencies



- Asset  $A_1$  depends on asset  $A_2$  if  $A_1$  cannot exist without some functionality of  $A_2$
- If  $A_2$  changes, so must  $A_1$ .
- Explicit in programs (module systems) and software projects (gradle)
- Used for modularization, *impact analysis*, *defect analysis*

## Example

Example RML:

```
1 roles:  
2   sources:  
3     - access: 'users.csv'  
4     referenceFormulation: csv  
5   s: dep:$role  
6   po:  
7     - [a, dep:Role]  
8     - [dep:roleName, $(role)]
```

Example SPARQL:

```
1 SELECT * {  
2  
3   ?x a dep:User;  
4   dep:name ?name;  
5   dep:hasRole [dep:roleName ?roleN].  
6  
7   FILTER (?roleN = "Admin")  
8 }
```

- Query depends on data output of engine driven by RML mapping
- Defect occurs, if we change URIs in the RML, but no tool can detect it!

# Semantic Assets

## Challenges

- Tools have no formal semantics, many domain-specific tools
- No explicit references
- Manual vs. derived assets

## Internal and External Semantic Assets

- An internal semantic asset is a mapping, a graph shape or a graph query.
- An external semantic asset is input data files, ontology axioms or source code operating on the final graph

We consider mostly RML mappings, not, e.g., python mappings

## External Dependencies

- A mapping  $M$  depends on a data file  $D$ , if  $D$  is input to  $M$
- A mapping  $M$  depends on an axiom  $X$  if  $M$  is generated from  $X$
- A program  $P$  depends on a semantic asset  $A$ , if  $A$  occurs within  $P$

```
1 roles:  
2   sources:  
3     - access: 'users.csv'  
4       referenceFormulation: csv  
5   s: dep:$role  
6   po:  
7     - [a, dep:Role]  
8     - [dep:roleName, $(role)]
```

## Internal Dependencies

- Partial order  $\preceq$  is the order of execution in the pipeline
- Library L is used to remove dependencies due to rdf:type etc.


Let  $L \subseteq \text{URI}$ . A semantic asset  $A_1$  depends on another semantic asset  $A_2$  if either

1.  $A_1$  refers to  $A_2$  explicitly, or
2. (2a)  $A_1 \preceq A_2$ , and (2b) there is some  $\text{uri} \in L$  that occurs in both  $A_1$  and  $A_2$ .

```
1 roles:  
2 ...  
3 s: dep:$role  
4 po:  
5 - [a, dep:Role]  
6 - [dep:roleName, $(role)]
```

```
1 SELECT * {  
2 ?x a dep:User;  
3 dep:name ?name;  
4 dep:hasRole [dep:roleName ?roleN].  
5 FILTER (?roleN = "Admin")  
6 }
```

## Example of a Dependency Graph



## Case Study: Teaching Ontology [SWJ, under review]

- 3 CSV files, 11 RML mappings, 19 SHACL shapes, 8 SPARQL Queries
- Fully automatic
- Found two bugs

### Bug 1: One query without dependencies

- Accesses data using a specific URI, but the mapping was commented out.
- Maintenance bug: Corresponds to an empty test for software.

### Bug 2: One shape without dependencies

- Change of URI prefix not propagated between dependencies.
- coursesonto:Lecturer vs. a local URI from the developer
- Maintenance bug
- Undetected because shape validation does not fail!




## Conclusion

---

ITU




# Conclusion

- The Semantic Web relies on software quality
- First steps towards investigating the field from this perspective
- Big challenges on the horizon: modularity and lack of formal semantics



# Conclusion

- The Semantic Web relies on software quality
- First steps towards investigating the field from this perspective
- Big challenges on the horizon: modularity and lack of formal semantics



Thank you for your attention